

ExecHQ Pilot Product: Functional Scope and User Journey

Product intent

ExecHQ is a private career advisor for senior managers through SVPs that turns a stated career direction into an actionable plan through concrete work, remembering what happened, and improving the next recommendation.

The pilot should validate one closed loop:

1. User states a direction or desired outcome.
2. ExecHQ recommends a templated plan with clear benchmarks and a high-confidence next action.
3. Active Landscape identifies what to do next and why.
4. The Toolbox helps the user produce a usable artifact.
5. The Loop captures use, outcomes, and feedback.
6. ExecHQ updates the record, tweaks the plan as needed, and recommends what to do next.

The Daily Intelligence Briefing is a separate daily reading product. It can feed the execution loop when relevant, but it is not the Plan's task list or execution surface.

The pilot is not a comprehensive presence audit, comparative benchmark, or always-on coaching product. It tests whether private, context-aware execution support creates immediate and continuing value.

Product principles and guardrails

- Private by default: users sign in with a personal email. Employer, cohort, and individual activity are never exposed to other users or the sponsoring organization.
- Direction, not title: users may state a role, desired scope, aspiration, or challenge. A precise target title is optional.
- Value before profile completion: only authentication and career direction are required before the first useful output.

- Agentic, not merely advisory: recommendations route into drafting, editing, exporting, and tracking real work.
- Transparent and non-gamified: show actions, artifacts, statuses, outcomes, and qualitative trends; do not show a comparative score.
- Pull over push: optional email can bring users back to something valuable. No streaks, leaderboards, or social mechanics.
- User control: every artifact is editable, exportable, and user-owned. Recommendations can be accepted, declined, deferred, or revised.

Pilot scope at a glance

Capability	Pilot decision	Rationale
Onboarding	Personal email and required career direction; optional refinement conversation and optional LinkedIn connection/import	Delivers value quickly. LinkedIn import is never required for first value.
First win	A recommended Toolbox action and draftable artifact	Proves ExecHQ does work with the user, not just planning.
Plan	Five suggested plan templates plus a guided custom-plan wizard (custom may be post mvp)	Provides a credible starting point while preserving flexibility.
Action Landscape	A dynamic, prioritized set of live next actions derived from the Plan	Is the concrete mechanism for advancing the plan
Toolbox	Positioning Builder, Pitch Builder, Situation Brief, Thought Leadership Builder	Each produces a useful artifact and enters the Loop.
Loop	Artifact status, user-reported use, outcomes, and recommendation feedback	Creates an observable record without scraping.
Daily Intelligence Briefing	A standalone daily experience with three curated reads, user-selected publisher preferences, and optional daily email	Creates a distinct reason to return and can route into Thought Leadership Builder.
Plan progress	Private activity record and qualitative activity picture (embedded in the Plan). Made up	Avoids scoring while preserving truth and directionality; shown in

	of the Signal Picture and Momentum Indicator	rolling 7-, 30-, and 90-day windows.
Monthly and quarterly reviews	Post-pilot / post-MVP	Requires enough longitudinal activity to be useful and credible.

Core user data model

Required at account creation

Input	Purpose	Notes
Personal email	Authentication, account recovery, optional email delivery	Enforce personal-domain policy for enterprise invitations.
Career direction	Drives the first recommendation and initial plan selection	Open text, such as “I want to lead a larger marketing organization” or “I need more influence in my current role.”

Optional context enrichers

These are offered after the first artifact is available, never required to receive it.

Input	Purpose	Notes
Refinement answers	Improve plan and artifact specificity	At most three high-value questions during first run, with a visible skip path.
LinkedIn connection/import	Improve user-approved career context, draft quality, and recommendation relevance	Explicit permission required. May be offered in onboarding or later as an Activity Landscape rec.
Signal Background	Maintain a private record of current context and self-reported signals	May include recent work, accomplishments, visibility activity, priorities, and changes.

Website connection	Improve user-approved publishing, thought leadership and presence context, as well as give context to voice and tone for writing	Explicit permission required. May be offered in onboarding or later as an Activity Landscape rec.
Weekly reflection	Refresh Signal Background through a lightweight journal exercise	A recommended context-building action, not a profile-completion gate.
Publisher preferences	Tailor the Daily Intelligence Briefing	User-controlled in Settings.

System-created records

- Direction statement and refinements
- Selected plan, roadmap, and plan rationale
- Active Landscape shown, accepted, declined, deferred, completed, and decline reasons
- Artifact drafts, versions, exports, and lifecycle status
- Outcome updates and user feedback
- Plan progress history by Plan area and rolling 7/30/90-day window
 - Signal Picture
 - Momentum Indicator
- Daily Intelligence Briefing engagement and resulting actions
- Signal Background and weekly reflections
- LinkedIn/Personal Website connection status, imported context, and explicit permission scope

Feature requirements

1. The Direction: onboarding and first win

When I want career momentum but do not have a precise next title, help me turn a broad direction into a useful first move without requiring a long profile.

User inputs

- Personal email and authentication

- Required open-text direction
- Optional choice to answer up to three follow-up questions
- Optional LinkedIn/Personal Website connection/import

System behavior

1. Present the privacy promise before collecting career information.
2. Ask: “What would you like to move toward?”* Validate both role-based and direction-based answers.
3. Interpret the direction to identify a likely need: positioning, visibility, influence, preparation for a specific moment, or career exploration.
4. Offer up to three high-value refinement questions with a visible “Skip and show my first step” path.
5. If skipped, use conservative defaults and state the assumption plainly.
6. Select one high-confidence immediate Toolbox action.
7. Open the relevant Toolbox flow, pre-populated with available context.

Immediate-value rule

A user who skips refinement must still reach a completable artifact, not an abstract plan. The default first win is normally Positioning Builder because it is broadly useful and can be refined later.

Example:

Based on your goal of leading a larger marketing organization, start with a concise leadership narrative. You can use it in a promotion conversation, executive introduction, or external opportunity. We can make this more specific later.

Outputs

- One-sentence interpreted direction, editable by the user
- One recommended first action with a “Why this, why now” rationale (an Active Landscape recommendation)
- A pre-populated Toolbox draft
- An invitation to refine the Plan after the artifact is saved or exported

Acceptance criteria

- A user can reach a first editable artifact using only email and direction.
- A user can skip all refinement questions without a blocked or degraded experience.
- The system makes no unsupported claims about history, employer, or performance.
- LinkedIn/Website connection/import is clearly optional and is not required for the first win.

2. The Plan: direction, roadmap, and live priorities

Job to be done

Help me see a credible path toward my direction while keeping immediate work focused and realistic.

Plan structure

Every Plan has two layers:

1. Roadmap: a long-term sequence of three to four stages over the user's stated or inferred time horizon.
2. Active Landscape: a deliberately small set of current priorities.

The active landscape contains:

- One to three short-term actions, generally achievable in 1-7 days
- One medium-term action, generally built over 8-30 days
- One long-term milestone, generally meaningful progress over a quarter or longer

Each action includes:

- Intended outcome
- Relevant plan area or signal
- Recommended timing and effort estimate
- Why this / why now / why you
- Clear completion definition
- Link to a Toolbox workflow or other in-product interaction, where applicable

Five sample plan templates

Plan	Best for	Example emphasis
Increase leadership scope	Strong performer seeking promotion or broader remit	Leadership narrative, executive-ready evidence, key conversations
Build executive presence	User needs greater visibility and a clearer point of view	Positioning, thought leadership, speaking or podcast pitching
Prepare for a career inflection point	Promotion, role change, review, board presentation, or negotiation ahead	Situation Brief, narrative, stakeholder strategy
Strengthen influence in the current organization	User intends to grow where they are	Strategic communication, executive presence, internal opportunity framing
Explore and clarify a next direction	User feels a ceiling but cannot name a role	Direction refinement, transferable narrative, low-risk exploratory actions

Plan selection and customization

- ExecHQ recommends one template based on onboarding inputs and explains why.
- The user may accept it, select another template, or choose “Build my own plan.”
- The custom-plan wizard gathers desired outcome, time horizon, current context, constraints, preferred growth areas, and available capacity.
- The wizard may use conversational follow-up, but it preserves progress and allows the user to exit with a draft plan.
- Users can accept, decline, defer, or edit each Active Landscape recommendation.

2.1 Active Landscape: what to do next, and why

Execution Recommendations (Active Landscape) are the Plan’s actionable next-step surface. They are not The Briefing.

Inputs

- Selected plan and active landscape
- Career direction and available context

- Artifact status and prior outcomes
- Stated capacity and preferences
- Recommendation acceptance, deferral, and decline feedback
- Signal Background, weekly reflections, and optional LinkedIn-imported context

Types of recommendation

- Artifact-producing action: create a Situation Brief, pitch, narrative, or thought piece
- Context-building action: connect/manually add LinkedIn, refresh Signal Background, complete a weekly reflection, or answer a targeted question
- Opportunity action: prepare for a specific meeting, moment, or audience
- Optional reading action: “Read today’s Daily Intelligence Briefing” when it is relevant to the Plan, and when appropriate create a directional Thought Leadership piece (LinkedIn Post)

Dynamic behavior

- Show no more than five active recommended actions at once.
- When a user declines an action, request an optional reason: not relevant, wrong timing, too much effort, uncomfortable channel, already done, or other.
- Immediately offer a relevant replacement from the ranked queue when available.
- When a user completes an action, offer to give the next appropriate recommendation without waiting for a weekly or monthly cadence boundary.
- Respect capacity: the user may accept the next action, defer it, or keep their current workload.
- Do not penalize a user for declining or deferring work.

Outputs

- A focused Active Landscape of live priorities
- Transparent rationale for each recommendation
- Feedback signals that improve future ranking

Acceptance criteria

- The user never sees a backlog disguised as a plan.
- Declined actions are replaced when suitable alternatives exist.
- Completed actions can trigger a next recommendation immediately.
- Plan recommendations are visibly distinct from the Daily Intelligence Briefing.

3. The Toolbox: artifact-producing workflows

All Toolbox flows must:

- Start from an Active Landscape recommendation, a Daily Intelligence Briefing item, or direct user initiation.
- Show intended use case, inputs needed, and relationship to the user's Plan.
- Gather only the context required for the artifact.
- Generate an editable first draft, not a final answer presented as fact.
- Support structured revisions and conversational follow-up.
- Save versions automatically.
- Support copy and export in pilot scope.
- Create a Loop record at draft creation.
- Offer optional context enrichers when useful, without gating access to the workflow.

3.1 Positioning Builder

Purpose

Create a reusable executive-level articulation of who the user is, the value they create, and the direction they are building toward.

Inputs

- Direction and Plan context
- Current role or function, if known
- Optional accomplishments, strengths, audience, and desired format
- User-provided source text, if available

Outputs

- Executive bio in short, medium, and long forms
- Leadership narrative: what I do, what I am known for, and what I am building toward
- Optional context-specific introduction or promotion-conversation opening
- Recommended next use of the artifact

User actions

- Edit directly
- Request revisions for voice, specificity, length, audience, or emphasis
- Copy or export
- Mark used and report outcome

3.2 Pitch Builder

Purpose

Create targeted, sendable pitches for speaking opportunities, podcast appearances, and press angles.

Inputs

- Chosen pitch type: speaking, podcast, or press
- Target outlet, event, or show, if known
- Career direction and Plan context
- User expertise and proof points
- Topical angle and audience
- Supporting links or prior materials

Outputs

- Tailored pitch draft
- Two to three subject-line or headline options
- Short and expanded versions where relevant
- Concise angle and “why now” statement
- Follow-up draft

User actions

- Edit directly
- Request revisions for tone, specificity, or length
- Copy or export
- Mark used or sent
- Report outcome

3.3 Situation Brief

Purpose

Prepare the user for a consequential review, board meeting, negotiation, executive conversation, or similar moment.

Inputs

- Situation type and date, if known
- Objective and desired decision or outcome
- Audience and stakeholder context
- Relevant evidence, risks, questions, or source material

Outputs

- One-page brief with objective, context, key messages, supporting evidence, likely questions, and requested decision or next step
- Opening statement or talk track
- Follow-up email draft when relevant

User actions

- Edit, revise, copy, or export
- Mark used
- Record whether the conversation created a next step or outcome

3.4 Thought Leadership Builder

Purpose

Turn the user's original idea, a Plan recommendation, or a Daily Intelligence Briefing item into an editable LinkedIn post or longer-form thought piece/article.

Inputs

- Starting point: original user idea, approved Active Landscape recommendation, or Daily Intelligence Briefing item
- Desired format: LinkedIn post, short article, or longer-form article
- User point of view, experience, evidence, and intended audience
- Direction and Plan context
- Desired voice, length, and call to action

Outputs

- Editable LinkedIn post draft
- Optional alternate hook and headline options
- Longer-form article or thought-piece draft when selected
- Suggested evidence gaps or questions for the user to address
- Recommended next use and follow-up action

User actions

- Edit directly
- Request revisions for voice, specificity, length, or audience
- Copy or export
- Mark published or used
- Report self-reported engagement or other outcomes

Guardrails

- Never publish automatically. The user authorizes and performs any external posting.
- Clearly distinguish user-provided evidence from generated framing.
- Do not infer or export private Signal Background information without explicit, scoped permission.

- Make provenance visible when a draft derives from LinkedIn-imported context or a Daily Intelligence Briefing item.

4. The Loop: action and outcome record

Job to be done

Help me keep track of career-building work I actually did, what happened next, and how that should change my Plan.

Artifact lifecycle

State	Meaning	User/system action
Drafted	Artifact created and saved	System records artifact type and linked Active Landscape recommendation, if applicable.
In progress	User is editing or gathering inputs	User may revise, pause, or abandon.
Ready	User indicates the artifact is usable	System prompts for intended use and timing.
Used / sent / published	User confirms it was used, sent, delivered, or published	System records date and optional channel or target.
Waiting	A response or event is pending	System schedules an appropriate follow-up prompt.
Outcome recorded	User shares result	System captures outcome type, confidence, and learning; updates Plan progress and refreshes the Active Landscape.
Closed	No further action needed	System updates Plan context and activity record.

Outcome inputs

- Outcome type: positive, neutral, negative, no response yet, or no longer relevant
- Optional detail: what happened, what the user learned, and what they want to do next
- Optional evidence or notes stored privately

System behavior

- Follow up only at contextually appropriate moments, with email as an opt-in channel.
- Summarize learning without overstating causality.
- Update future Active Landscape recommendations based on outcomes, acceptance and decline patterns, capacity, stated preferences, Signal Background, and optional LinkedIn-imported context.
- Preserve an accessible, exportable Plan progress record.
- When an action is completed or an outcome is recorded, offer the next appropriate Active Landscape recommendation without waiting for a fixed cadence.

Pilot measurement language

Use “completed action” and “Plan progress,” not a score or gamified momentum indicator.

Example: “You completed two planned actions this month and have one artifact awaiting an outcome.”

5. The Plan: progress and momentum

The Plan is the user’s single place to see direction, roadmap, Active Landscape, and progress over time. It provides a private, factual view of activity tied directly to the Plan rather than a separate Progress feature.

Show:

- Active Plan roadmap and milestones
- Active Landscape and current live priorities
- Plan progress by Plan area in rolling 7-day, 30-day, and 90-day windows (private, factual; never a comparative score) — each Signal Picture item reports evidence of movement against a specific Plan area, Plan action, or milestone
- Artifacts and outcomes linked to Plan actions
- A concise Plan progress narrative: what changed and what next, explicitly tied to Plan areas and how evidence refreshes the Active Landscape

Signal Picture (embedded in the Plan): clearly labeled platform-observed and user-reported activity across available presence signals, shown in 7-, 30-, and 90-day windows. Examples: 7 days — drafted one post, sent one podcast pitch, read three Daily Intelligence Briefing articles; 30 days — published two posts, submitted one speaking

abstract, received one press mention; 90 days — published content activity is increasing, speaking activity is unchanged, press presence remains light. Individual items are factual; longer windows show the user's trajectory, not a comparison with others.

Momentum (in the Plan): a transparent execution trend, not a score. 7-day view: actions completed, artifacts created or used, and outcomes updated. 30-day view: consistency of follow-through and progress against the Plan. 90-day view: directional momentum labeled as building, steady, or needs attention. Underlying activity is visible; users are never reduced to a composite score, rank, or behavioral grade.

What changed / what next: a concise narrative grounded in the 7/30/90 windows and linked to the Active Landscape. Example: "Over the past 30 days, you published twice and completed three Active Landscape actions. Over 90 days, your content activity is building, while speaking outreach is unchanged. Your next best move: send the conference pitch you drafted."

Do not show:

- Composite Growth Index
- Rank, percentile, leaderboard, streak, or prescriptive score
- Claims that activity alone caused a promotion, visibility outcome, or career result

Post-MVP measurement expansion

Once meaningful activity exists, add a monthly review and later a quarterly Executive Presence Audit.

Surface	Timing	Inputs	Outputs
Monthly Progress Review (Plan)	After 30 days of meaningful activity	Completed actions, artifact states, user outcomes, optional self-reported signal activity	Narrative of work completed, trends, gaps, and next-month focus
Quarterly Executive Presence Audit	After one full quarter	30/90-day activity, outcomes, optional verified or self-reported external signals, Plan themes	Private qualitative report, signal picture, reflection, and revised roadmap

- Platform-observed (in Plan): artifacts created, actions completed, and outcomes entered in ExecHQ
- User-reported (in Plan): publishing, appearances, mentions, and other external activity

- Future integrated: signals from permitted, privacy-safe sources (when incorporated into Plan progress)

6. The Daily Intelligence Briefing: three reads worth your time

Give me a small, tailored daily reading experience that is worth opening, rather than an endless professional-content feed.

The Daily Intelligence Briefing is a standalone site destination and optional daily email. It is not the user's Plan, task list, or Active Landscape surface.

Inputs

- Career direction and selected Plan
- Publisher preferences selected in Settings
- Stated interests and optional Signal Background context
- Limited curation and rules-based relevance in pilot scope

Outputs

Each daily session includes:

- Three curated articles from the user's chosen or approved publisher set
- A concise explanation of why each item matters to the user's direction or interests
- An option to open Thought Leadership Builder and draft an editable LinkedIn post or article response

User actions

- Read or save an item
- Indicate read/not read
- Provide reaction or relevance feedback
- Choose whether to generate or refine a post
- Adjust publisher preferences and daily email delivery in Settings

Relationship to the execution loop

The Briefing may become an input to the execution loop only when ExechQ explicitly presents "Read today's Daily Intelligence Briefing" as an Active Landscape

recommendation. This is optional and contextual. The Briefing is otherwise a distinct reading experience.

Pilot boundary

Use curated and rules-based relevance with limited personalization. Do not position it as comprehensive market intelligence. No automatic publishing occurs.

End-to-end user journey narrative

Persona

Maya is a senior marketing leader at a large organization. She is successful but feels peers with clearer executive narratives and broader visibility are moving ahead. She wants greater scope but is not ready to declare an active job search.

Day 1: value before interrogation

Maya creates an account with her personal email. Before entering career information, ExecHQ explains that her employer and colleagues cannot see her activity.

She enters: "I want to move from running campaigns to leading a broader marketing organization." ExecHQ offers three optional questions. Maya is between meetings and selects "Show me a first step."

ExecHQ responds: "A clear leadership narrative will make this direction usable in conversations about scope, promotions, and opportunities. Let's draft one." Maya enters a few proof points and receives an editable executive narrative, short bio, and opening for a future scope conversation. She revises the language and saves the artifact.

ExecHQ then offers an optional LinkedIn connection to improve future draft context. Maya skips it. She also sees a low-effort Active Landscape Rec to add three recent accomplishments to her Signal Background later; it is not required.

ExecHQ recommends the Increase Leadership Scope Plan. It explains that Maya's direction suggests a need to clarify her leadership proposition, build evidence for broader scope, and prepare for the conversations that determine it.

Her Execution Recommendations include:

- Finalize her leadership narrative.
- Create a Situation Brief for her upcoming manager check-in.
- Identify one cross-functional initiative where she can demonstrate broader scope.

The Plan also includes one medium-term action: prepare a point of view for an internal leadership forum; and one long-term milestone: establish a documented case for expanded scope within the next quarter.

Maya accepts the first two actions and declines the cross-functional initiative because the organization is in a planning cycle. She selects “Wrong timing.” ExecHQ immediately replaces it with a relevant alternative: prepare a stakeholder map for the planning work already on her calendar.

Week 1: produce and use

Before her manager check-in, Maya opens the Situation Brief from an Active Landscape recommendation. She supplies the meeting purpose, desired outcome, and recent accomplishments. ExecHQ produces a one-page brief and opening statement. Maya edits it, uses it, and marks it “Used.”

Two days later, ExecHQ asks whether the conversation created a useful next step. Maya reports that her manager asked her to lead a cross-functional planning workstream. She records that she needs to demonstrate executive-level stakeholder management.

ExecHQ updates its context and offers a new Active Landscape recommendation immediately: build a stakeholder message map for the workstream. It explains why: Maya said the assignment requires cross-functional influence.

Week 2: return through daily relevance

Maya opens her Daily Intelligence Briefing, a separate destination containing three articles from publishers she chose in Settings. One item is relevant to her planning workstream. It explains why it matters and offers: “Turn your perspective into a LinkedIn post.”

Maya opens Thought Leadership Builder, adds her point of view and an example from her experience, and produces an editable post. She saves it as a draft; ExecHQ does not publish it for her.

The Briefing remains a daily reading experience. The Plan can later recommend reading it when useful, but the articles themselves are not Maya’s execution task list.

Month 1 and beyond: continuity becomes value

Maya’s dashboard shows factual Plan progress: artifacts created, work used, outcomes recorded, and current Active Landscape recommendations. It does not score or compare her.

At the end of the month, an optional progress review (in the Plan) says that she completed planned actions, created reusable artifacts, and converted a manager

conversation into a broader-scope assignment. It identifies a practical next gap: preparing a consistent point of view for the stakeholders she now needs to influence.

Over time, Maya's record becomes valuable because it retains the work, decisions, outcomes, and narrative she would otherwise need to reconstruct from memory.

Pilot success measures

Question	Evidence to collect
Does the first win land?	Percentage reaching an editable artifact; time to first artifact; perceived relevance and usefulness in research sessions
Does the loop close?	Artifact-to-used conversion; used-to-outcome-recorded conversion; time to outcome update
Does personalization feel credible with limited inputs?	First-recommendation acceptance; decline reasons; qualitative relevance feedback; optional LinkedIn-import uptake and perceived value
Do Active Landscape recommendations feel useful and responsive?	Completion, decline, replacement, and next-action acceptance rates; feedback on workload and timing
Does the Daily Intelligence Briefing earn standalone return use?	Daily open/read rates, saves, relevance feedback, Thought Leadership Builder starts, and user-stated value
Does the Plan feel focused rather than generic?	Plan adoption, active-action completion, and user feedback on clarity and workload; Plan-progress usefulness
Does the record feel worth keeping?	Return rate, prior artifact/outcome revisits, and stated intent to continue
Is candor useful rather than discouraging?	Feedback on narrative summaries and recommendations; abandonment following direct feedback
Do context-building recommendations help?	Uptake and downstream impact of LinkedIn/Personal Website connection/import, Signal Background refresh, and weekly reflection

Explicitly out of scope for the pilot

- Numerical Growth Index, ranking, percentile, and leaderboard mechanics
- Required LinkedIn connection, scraping, or inferred external activity
- Automated sending, publishing, or outreach on the user's behalf
- Full opportunity database, recruiting marketplace, or peer review
- Quarterly report generation before enough longitudinal activity exists
- Broad enterprise reporting beyond anonymized pilot participation and product-health signals
- Additional Toolbox workflows and signals

Product decisions to lock during design and engineering

1. Approve Positioning Builder as the default first-win workflow when a user skips refinement.
2. Confirm the five Plan templates and custom-plan wizard inputs.
3. Confirm Active Landscape as the Plan's dynamic next-step surface, distinct from the Daily Intelligence Briefing.
4. Confirm dynamic behavior: replace declined recommendations when possible and offer the next appropriate recommendation after completion.
5. Confirm the four pilot Toolbox workflows and their exact output formats.
6. Define status semantics for used, sent, published, outcome recorded, and closed.
7. Confirm Daily Intelligence Briefing publisher preferences, daily-email delivery, article sourcing, and feedback prompts.
8. Define observed versus user-reported activity and prohibit blending the two without clear labels; embed Plan progress in the Plan row.
9. Confirm privacy architecture, personal-email enforcement, employer non-visibility, export/deletion expectations, and enterprise reporting limits.

Alpha: Prove the core loop

Goal: Can a user get a useful artifact, act on it, and see that ExecHQ remembers the result?

Include

1. Lightweight onboarding
 - Personal email
 - Open-text career direction
 - Optional 1–3 follow-up questions
 - No LinkedIn import required for alpha
2. First win
 - Default into Positioning Builder
 - Produce an editable executive bio / career narrative
 - Explain “why this, why now”
3. A minimal Plan
 - One stated direction
 - One current Plan stage
 - 2–3 suggested actions, rather than a full staged roadmap
4. Active Landscape
 - One prioritized next action at a time
 - Accept, decline, or defer
 - Declining yields a replacement
 - Completing an action yields the next appropriate action
5. Two Toolbox tools
 - Positioning Builder
 - Thought Leadership Builder
 - These support the first win and a repeatable follow-up action without needing complex external integrations.
6. Minimal Loop

-
- Drafted → used/published → outcome recorded
- User-reported outcomes only
- A simple history of artifacts and completions

7. Plan progress

- A private “since you started” record
- Completed actions, artifacts created, artifacts used, outcomes recorded
- Do not build 7/30/90 trend views until enough history exists.

Exclude from alpha

- LinkedIn import and other integrations
- Daily Intelligence Briefing
- Pitch Builder and Situation Brief
- Full multi-stage Plan
- Signal Picture across six channels
- 7/30/90 progress trends
- Enterprise dashboard
- Automated follow-ups, notifications, or outcome detection
- Quarterly Audit/reporting

Beta / Pilot: Prove repeat value

Add the parts that make ExecHQ feel like an ongoing product:

- Full Plan with roadmap, stages, and milestones
- Active Landscape with up to five live actions
- All four Toolbox tools
- Daily Intelligence Briefing
- Optional LinkedIn connection/import
- Signal Background and lightweight weekly reflection

- 7-, 30-, and 90-day Plan progress views
- Plan-linked Signal Picture, Momentum, and “what changed / what next”
- Optional daily briefing email
- Moderated research and instrumentation

Full V1: Prove the product can scale

Add after pilot evidence supports the core behavior:

- Broader connected-account support, subject to legal review
- More complete Signal Picture coverage
- Automated or assisted activity capture
- Quarterly Executive Presence Audit / Report
- Additional Toolbox tools
- More robust recommendation intelligence based on validated patterns

The key is that alpha should not try to prove “career progress measurement.” It should prove the prerequisite: users will let ExecHQ turn their direction into work, use that work, and return to close the loop.