


ExecHQ

PRODUCT STRATEGY, V1 SCOPE & POST-LAUNCH ROADMAP

General Outline

1. What We Heard
2. The Challenge
3. The Audience
4. The North Star
5. V1: The Pilot Product
6. The Roadmap
7. Next Steps



What We Heard

Privacy is settled. The experience is still open.

Nothing a user does can be visible to their employer, their network, or another user. That is the one fixed point.

Major features are open: described as up for interpretation, and could be deleted or enhanced entirely.

The index is provisional: by no means the definitive benchmark, and no authoritative version of the formula exists.

Onboarding is acknowledged as broken: what it is and what it should be are two different things.

The UX is ours: the first-run experience is open to all recommendations.

“ To me, the privacy is absolute. That is the core. ”

“ The rest of the stuff is cool, but could it be either deleted or enhanced? Sure. ”

We have no shortage of good ideas, but no hierarchy yet

The prototype is a catalogue of ideas, nothing in it is ranked, and several items do the same job.

Overlapping

- Growth Engine
- Growth Path
- Growth Advisor
- Confidant

Four names for one idea.
Flagged for consolidation.

Measurement

- Growth Index
- Momentum Score
- Presence Simulator
- Score History

Two competing scores,
and no stated relationship
between them.

Content & Intel

- Daily Briefing
- Why Now
- Weekly Pulse
- LinkedIn generator

The most defended group,
and the closest to working
today.

Enterprise

- Leadership Bench
- Flight Risk
- Development ROI

To be defined

One dashboard is required for Deloitte. What it contains is undefined. See slide 32.

The score already drove a real user away

A beta tester who does not use LinkedIn opened the product, saw their number, and did not come back.

20%

of the Growth Index is LinkedIn presence, the single heaviest signal in the formula.

A score is a verdict. It arrives before any value does, and the user has no way to argue with it.

It measures the wrong distance. Hitting 90 guarantees nothing, so the number and the goal come apart.

Its real audience was never the user. The index existed partly to flag high-value people to partners.

“ They felt like they were a “loser” immediately because their score was so low. ”

“ I think it can be construed more negatively than positively. ”

The replacement arrived overnight

A full proposal came in Saturday morning to replace the Growth Index with a qualitative audit. It answers the same question without a number.

The Signal Picture: what the user actually did across the six channels over 30 and 90 days, presented as facts rather than converted into a score.

The Momentum Indicator: a behavioral measure that advances only through completed tasks. Binary, transparent, and not gameable.

The Quarterly Report: a narrative document the user keeps, combining the picture, the completion record, and the themes from their sessions.

What it removes: the conversion of activity into a number. The categories stay, the scoring goes.

What it costs: the Audit as written depends on roughly eight tools we have not scoped. The model is right; the dependency list is not V1.

“ The number requires an external reference point to be meaningful, and that reference point is either arbitrary or unavailable. ”

“ Career momentum is not produced by having a LinkedIn score of 68. ”

The instincts are right: private, honest, agentic

Four convictions held across all three sessions. These are the constraints we designed against.

Private	No employer, no network, no audience. The platform knows you well; nobody else knows anything.
Not gamified	No streaks, no leaderboards. Gamified is on the brand document's list of words to never be.
Pull, not push	Built for career inflection points. Users should seek it out, not be nagged toward it.
Agentic	Not just tell me to send three insights, but help me draft them.

“ I don't want to gamify this. To me that's hokey, and there's very little value there. ”

“ Take the friction out of the experience for the user, and that's where the clear value proposition lies. ”

The six categories are right, but they are not the whole picture

The framework for external presence is defensible and stays. What it cannot do alone is serve a manager and an SVP with the same formula.



What has to sit alongside them

- Professional development completed
 - Mentorship, given or received
 - Internal scope and influence
 - Team and org leadership
 - Communication and presentation
 - Anything a manager can realistically move in ninety days
- These are not V1. They are what the framework grows into once there are tools that produce them. All six current signals ask the same question: is the outside world looking at you? None of them move for a manager inside a large firm in a single quarter.

The six external signals, as weighted today

"A manager or a director, the chances of them becoming a thought leader are very, very small versus an SVP." Keeping the categories is right. Making them the only inputs is what breaks the bottom half of the audience.

“

What is truly special about this thing that cannot be copied, and can continue to provide momentum and acceleration over time, regardless of the inflection point that I come back from or to?

”

This deck is our answer.

The Challenge

What a prompt can't do is come back tomorrow

We gave Claude a short description of a VP of growth marketing who wants to become a CMO. It returned a document as specific as the app, in five minutes.

A 30-month plan in three phases, each pairing an internal scope milestone with an external visibility one.

Named outlets with their actual submission processes and contributor requirements.

Named podcasts and hosts, ordered by which are realistic to pitch at VP level.

A worked 12-week schedule and a weekly cadence totalling five to six hours.

Named search firms to build relationships with before going to market.

“

More informative, action oriented and contextual than the app. But how do we beat that? I gave it a four sentence prompt.

”

Two of the nine proposed Toolbox tools already exist inside that one output. The differentiator cannot be just the output.

Executive coaching today, recruiting tomorrow


The product complements a lot of things. What it displaces first is the executive coach, and the reason is cost, cadence, and the absence of a record.

Executive coaching	Expensive, slow, no measurement layer, and only as good as the individual coach.	Day One Target
LinkedIn	Public by architecture. Its business model requires visibility; ours requires the opposite.	Complementary
LinkedIn audit tools	Commoditizing fast, and solving a narrower problem: how your profile looks.	Complementary
Recruiters and search	Sees you only when a role is open, and works for the employer, not for you.	Future Target

Coaching is what a user chooses between today. Recruiting is the data play that opens once enough people have used it.

A summary is an input, never a decision

We are all using AI on this project, us included. Here is how we would suggest we work from here.



Summaries generalize

That is fine for notes. It is a problem for requirements, where the details are the whole point. Reading the summary and deciding what matters is the part a person has to do.

The Audience

V1 is built for one person, and we know her name

The senior manager who is doing well and can feel the ceiling

- 35 to 50, manager through SVP
- Media, marketing, communications, and tech
- Inside a large organization, signing up on a personal email
- Accomplished, not struggling
- Has watched a peer get promoted and cannot say why

Not in V1

C-suite (crowded, and a competitor is already aimed at it) · Early career (that is EmergeHQ) · Career changers into unrelated fields · The genuinely lost, who need a different first-run experience

“ Deloitte believes in empowering individuals within their organization, whether they stay with them or not. ”

“ Manager to VP or SVP. That's a little bit more condensed. ”

Requiring a job title will break onboarding

The current design opens by asking for an aspired role. Everything downstream depends on that answer.

THE ASSUMPTION

For the most part, people have an idea of where they want to go.

THE BRAND DOCUMENT

The secondary user is the executive who is genuinely stuck and knows it. This user is vulnerable.

Both cannot drive the same first screen

Our recommendation for V1 is to ask for a direction, and accept a title. "CMO in three years" is a perfectly good direction if the user has one, and we should take it and run with it. What breaks is requiring a title from someone who cannot name one. "I want to run marketing" and "I want out of agency life" have to be equally valid answers in the same open field. A direction is enough to generate the first result, and it sharpens over time.

“

75% of my friends are calling me because they are unhappy with their jobs and thinking about a change, but no one can name what they want to do next other than make more money.

”

“

Work life balance is the main driver.

”

The North Star

Senior professionals have always had to navigate their most important career decisions without the one resource those decisions actually require — a completely private, rigorously honest, and deeply specific intelligence layer that is accountable only to them.

ExecHQ is the private career growth and momentum platform — combining personalized coaching, advisory guidance, curated daily intelligence and tools into one resource built entirely around each individual's goals, signals, and desired path forward.

Anyone can get advice in two minutes. Nobody else keeps the record, produces the artifact, or risks the relationship by being honest.

Three things a general-purpose AI will never do

Each maps to a pillar already written into the brand document. We are not inventing a position; we are choosing which one to build on.

1

Continuity

It remembers what you did, not what you asked.

A chat session starts over every time. ExecHQ holds a record of actions taken and outcomes, and gets more specific the longer it runs.

2

Execution

It produces the thing, then asks what happened.

Not "pitch this podcast" but a drafted pitch, marked sent, followed up. The plan and the doing are the same surface.

3

Candor

It tells you what your coach will not.

A coach you pay is invested in your comfort. A chatbot agrees with you. ExecHQ is the only one that can afford to be blunt.

The promise is about who can see it, not what we know

These two words get used interchangeably and they point at opposite products. The distinction decides how onboarding works and how good the first recommendation can be.

IF WE MEANT ANONYMITY

- The platform knows nothing it is not handed
- Every signal typed by hand, then typed again to update it
- Recommendations built on four fields
- A first win any prompt could have produced

BECAUSE WE MEAN PRIVACY

- Connect the accounts they already have
- Know their history, their writing, their presence on day one
- Recommendations grounded in what they have actually done
- A first win nobody else could have produced

The fixed point has not moved

Nothing a user does is visible to their employer, their network, or another user. That is a promise about exposure. It was never a limit on what the platform is allowed to know about the person it is coaching.

“

We can keep your information private, but we need information about you to get you the plan you need.

”


“

I have to remember to manually go back and update my stuff. It would be better if there was some kind of listener that auto-fills my updates.

”

One mechanic makes all three of them real

This is the single most important idea in the deck. It is what makes progress something we can verify rather than something the user asserts.



Because the artifact originates in the platform, the platform knows it exists

That is observable behavior whether or not integrations land and no terms-of-service exposure. It is also the only progress measure that is honest, because it tracks what the user actually did rather than what they typed into a form.

The moat gets deeper with every user

Candor and execution are available on day one with fifty users. The defensible asset arrives later, and only if we start collecting it now.

NOW 50 users	Candor + Execution	Honest guidance and real artifacts. Available immediately, and enough to differentiate in the pilot.
NEXT 1,000 users	Pattern data	What the first thousand records show about which actions actually precede advancement, fed back as better recommendations. Aggregate only, no identity, nothing exposed to another user.
LATER 10,000+ users	Aggregate outcomes	What actually worked for people like you. Longitudinal, proprietary, and impossible to copy without the same years of data.

Recommendation: during the pilot we run the research ourselves, in moderated sessions. Nothing a user does is ever visible to another user, in the pilot or in the product.

V1: The Pilot Product

Onboarding: The first ninety seconds decide everything that follows

Aspired direction, the privacy promise, and one usable result before we ask for anything else.

What it is

- Three steps: email, direction, first win
- Privacy shown before a single field is requested
- An open field with guardrails. A title is a valid answer, just not a required one.
- Optional: up to three refinement questions (maybe LinkedIn connection/import) Skip is always available.
- First introduction to a quick win and toolbox item including the “Why this, why now” context.

Why it earns the slot

- The current onboarding asks for a lot and explains nothing
- The first win must be real output, not a tutorial. A tour is a promise; a result is proof
- Requiring a target title first filters out the user who needs this most
- This screen is where privacy either becomes believable or becomes marketing copy

“

It should just explain very succinctly what the value proposition is, rather than give us a bunch of information first.

”

“

User aspirations and future goals matter more than scraping historical job data.

”

Features list is consolidated and scaled down to five core pillars

Everything else is deferred, not deleted. The roadmap in Section 6 when we fold them back in.

- | | | |
|---|--|---|
| 1 | The Plan
One consolidated growth path | Because four overlapping features confused everyone, including us. |
| 2 | The Active Landscape
Personalized coaching that adjusts with you | Because something that works for you not with you gets lost. |
| 3 | The Toolbox
Tools that execute the follow through | Because advice without an artifact is what every LLM already does. |
| 4 | The Loop
Outcome tracking on every step | Because accountability keeps users coming back and it does not exist yet. |
| 5 | The Briefing
Daily intelligence, tied to you and the plan | Because it is a sticky way to pull or push and a good way to help users see value and stack wins. |

The Plan: Progress is a picture, not a number

One growth path replaces the Growth Engine, Growth Path, Growth Advisor, and Confidant, with the Presence Audit as its progress surface. Progress is shown through a private Signal Picture, execution momentum, and the record of what changed over time.

What it is

- A staged plan with a few live next actions, never a backlog of forty
- 3-4 stages toward the user's goal. Progress shown as the Signal Picture plus completed tasks, never a number
- Gives the user an enduring view of where they are headed.
- Signal expansion comes later. V1 tracks only signals a tool can actually move.
- A visible reason on every recommendation: why this, why now, why you

Why it earns the slot

- Four features doing one job was the clearest consolidation call of the sessions
- The Audit answers "where do I stand" without the comparison that drove a tester away
- Every observation ties to something the user did, which is what makes it actionable
- The Quarterly Report is the right idea on the wrong clock for a three-week pilot. See the deferred list

“

The Audit produces a narrative. Every observation is connected to something the user did or did not do.

”

“

It is about the journey, not the destination. An executive coach isn't guaranteeing anybody a job.

”

The Active Landscape: Things change, the plan should change with them

As things change for better or worse in the real world, the ExecHQ LLM will work with the user to make adjustments each next step as needed.

What it is

- A personal coach working with the user to adjust for whatever life throws their way.
- A dynamic, prioritized set of up to 5 live actions derived from the Plan — not a backlog (Includes artifact work, context-building, or “Read today’s Daily Intelligence Briefing” when useful)
- A visible reason on every recommendation: why this, why now, why you
- Also an avenue for the user to ask questions, weigh ideas about what their goals should be.
- Feeds the Toolbox directly: read this, now draft the reaction

Why it earns the slot

- Each user will be on different timelines for executing their plan. We need a tool that understands that.
- A recommendation is not valid without evidence to support it. This feature explains why the move is relevant to the user and what it will do for them not only at this moment in time but in the future as well.

“

That can be adjusted at any time by the person, and the platform reacts to that.

”

“

There needs to be a kind of a reasoning behind why.

”

The Toolbox: The gap isn't knowing what to do. It's doing it.

Four tools, cut down from nine. Each produces something the user can physically use, and each moves a signal the plan is tracking.

What it is

- Pitch Builder: speaking abstracts, podcast pitches, press angles
- Positioning Builder: bio and career narrative in every length
- Situation Brief: one page before a review, a board meeting, a negotiation
- Thought Leadership Builder: editable LinkedIn posts and longer-form thought pieces or articles that translate the user's perspective into public-facing content
- Every output editable, exportable, and owned by the user
- Every output enters The Loop the moment it is created

Why it earns the slot

- These items are sendable, tied to a signal that moves, and used repeatedly. The other six fail at least one test
- Bio and Career Narrative were two proposals doing one job. Merged
- Signal Sprint Planner and Weekly Build Plan are The Plan, not separate tools
- Interview Intelligence, Decision Framework, and Reference Prep are strong, lower frequency, deferred to Q1

“ Can it not just say send these three unique insights, but let me help you now draft these three unique insights? ”

“ ExecHQ needs to become more indispensable to the user. Otherwise it's a nice-to-have. ”

The Loop: If the platform made it, the platform can measure it

Status tracking on every artifact the platform produces. This is where our thinking and the Presence Audit already agree, and it is the part of progress we can actually verify.

What it is

- Every artifact carries a state: drafted, sent, waiting, outcome
- The platform follows up and asks what happened
- Completed tasks are the input the Momentum Indicator runs on
- Outcomes feed the plan and change what is recommended next
- Completed actions, used artifacts, and reported outcomes update the user's private progress record.
- Works whether or not integrations land, because the platform made the artifact
- The record becomes the evidence for your own review, instead of reconstructing your wins from memory in December.

Why it earns the slot

- The Audit already calls completed tasks binary and not gameable. The Loop is the mechanism that makes that true
- Without it, progress is self-attested. The user types that they did the thing and the platform congratulates them.
- The other half of the Signal Picture should come from connected accounts, not a form.
- It is the difference between a document and a product. A plan you can ignore is a PDF
- It creates the longitudinal record that becomes the real moat at scale

“

Advances through completed Toolbox tasks. Cannot be gamed. It measures behavior, not outcome, and is entirely transparent.

”

“

The platform tracks them through user updates and third-party integrations if available.

”

The Briefing: A curated list beats an infinite feed

Each content piece is tied to the plan or the user, and carries the reason it matters to them this week.

What it is

- A daily collection of relevant content and possible reminder of the day's quick win, a weekly bigger move, a monthly long-term play.
- Curated against the user's plan and the execution recommendations, not general news
- Each item carries a one-line reason it is here
- Optional email or text nudge. No in-app notifications.
- Actions: Read, save, mark relevant/not relevant, and open Thought Leadership Builder to draft a LinkedIn post or article response.

Why it earns the slot

- The most defended feature across all three sessions, and closest to working today
- It is the pull mechanic, and the reason to open the product between inflection points
- The reason line is the difference between this and an RSS reader

“

It knows me. It's really refining things to my specific area of interest.

”

“

I'm not so keen on things popping up on your phone and it's swipe away, swipe away.

”

Five features ride on six things nobody will ever see

None of the five ship without these. They are not negotiable, and they are not visible.

Privacy architecture

No employer visible in any user-facing surface. Org membership exists only as an invite record, never as profile data.

Real authentication

Accounts replacing a beta where anyone with the URL can act as any user.

Personal email enforcement

Server-side validation blocking corporate domains on enterprise invitations.

Server-side AI calls

Every model call proxied through the backend. No key reaches the client.

Cost controls at scale

Per-user caching and tiered call handling, tested to 1,000 concurrent users.

Three responsive breakpoints

Mobile, tablet, and desktop tested independently. This audience has no tolerance for a clunky first two minutes.

Everything V1 does fits on one page

This is the scope at the level a decision can be made from, with what each surface does and what waits.

Surface	In V1	Waits
Onboarding	Direction in an open field, account connection, one usable output inside ninety seconds	Long profile forms, a required job title
The Plan/Active Landscape	One staged path. A few live actions across daily, weekly and monthly horizons, each with a stated reason	Four parallel growth features, any 0 to 100 score
Metrics	Signal Picture: factual 7- 30- and 90-day activity across available presence signals. Momentum: visible execution trend based on completed actions, artifacts used, and outcomes reported. No composite score or user comparison	The Quarterly Report as a document
The Toolbox	Pitch Builder, Positioning Builder, Situation Brief. Thought Leadership Builder. Every output editable, exportable, owned by the user	The other six proposed tools
The Loop	Drafted, sent, waiting, outcome on every artifact. Follow-up prompt, and the plan updates from the answer	Automatic outcome detection
The Briefing	Three curated reads from user-selected publishers; optional daily email; relevance feedback; optional Thought Leadership Builder handoff.	In-app push notifications, an infinite feed
Enterprise	One dashboard, cohort level only, no individual identity inside it	Flight Risk, Leadership Bench, Development ROI
Platform	Real accounts, personal email enforcement, server-side model calls, three responsive breakpoints	Anything that puts a key on the client

Everything in the Waits column has a date on the deferred list and in the roadmap.

The user is not the data entry layer

Connect what can be connected, ask for the rest a little at a time, and never ask twice for something we can already see.

1 CONNECT

OAuth into the accounts they already have. LinkedIn first, then whatever else clears legal review. Offered one click at a time, at the moment it obviously helps, not as a wall on the first screen.

2 INFER

Whatever the connection returns becomes the profile. Writing, speaking, presence, history. The user confirms and corrects rather than composes.

3 ASK

What is left gets asked progressively, a question or two at a time, in conversation, and only when the answer changes the next recommendation. Offer Signal Background and a lightweight weekly reflection as optional, user-controlled ways to keep context current.

Still open: which integrations are legally available, and through which provider. Dom & Tom are auditing this week.

“

People are used to having systems do the work for them. Having to do it manually is a step backwards.

”

“

There is a cleverness to teasing out people's data, and that is where engineering can help. But it all has to be purposeful.

”

The dashboard is a major part of the experience, now it needs definition

It is not out. It is required for Deloitte and undefined, which makes it the largest unscoped item in V1.

What is settled

- One dashboard, not several
- An enterprise surface, separate from the five user features
- No names inside it. Cohort membership only, never individual identity.

What is open

- What Deloitte actually wants. Everything we have is secondhand
- Whether it reports cohort activity, program health, or individual momentum
- Whether it ships for the 50-user pilot or first appears in December

Three outstanding questions

- 1 Who at Deloitte has seen this, and what were they promised?
- 2 Is the December audience buying the dashboard, or the user outcome?
- 3 What is the smallest version that satisfies the partner meeting?

“ We have no idea what Deloitte is thinking, other than what we are getting secondhand. ”

“ They do not know what they do not know. They will react to whatever we put in front of them, so we should decide what that is. ”

Nothing is cut, and everything has a return date

None of this is rejected. Each item has a reason and a date it comes back.

The Quarterly Report

Q1

The quarterly narrative document waits until sufficient history exists. The underlying Signal Picture and Momentum Tracking begin in V1.

Any numerical score

Revisit

Failed with a real tester, and withdrawn on 9/12. May persist internally for partner matching only.

Peer Review

Not planned

Introduces an audience, which contradicts the privacy promise. Revisit only if users ask for it.

Streaks, share cards, Momentum Feed

Not planned

Conflicts with the anti-gamification stance and the brand document.

Enterprise reporting suite

Q1

One minimal dashboard ships in V1. The full suite comes later.

Confidant as a paid tier

Q1

Single tier in V1. No gating until there is something proven to gate.

Six remaining Toolbox tools

Q1 to Q2

Strong ideas, lower frequency. Sequenced by what the pilot shows us.

IC versus management paths

Q2

Valid, and explicitly deferred past MVP on 9/11.

The Roadmap

Each phase answers a question the last one could not

Four phases. Every one of them exists to settle something specific before the next begins.

NOVEMBER

The Pilot

50 Deloitte users

- The five features
- Privacy architecture
- Moderated research sessions, weekly

Does the loop hold? Do people come back?

DECEMBER

The Partner Meeting

Scale to 1,000

- Fixes from pilot findings
- Load and performance at scale
- Terms of service and privacy policy live
- Demo storyline for the stage

Does it hold up in front of an audience?

Q1 2027

Depth

Post-pilot

- Pattern data from the first cohort
 - Remaining Toolbox tools
 - Confidant and a paid tier
 - Enterprise reporting beyond the V1 dashboard
- What will people pay for?

2027

The Moat

Mass market

- Aggregate outcome intelligence
- Enterprise reporting suite
- EmergeHQ on the shared engine
- Naming resolved

What can no one else copy?

Fifty users can only answer four questions

Three weeks with fifty people is not a lot of data. These are the only questions worth spending it on.

1

Does the first win land?

Does a new user reach a usable output without help, and do they believe it was made for them?

2

Does the loop close?

Do users actually send the artifacts, and come back to say what happened? If not, the differentiator is theoretical.

3

Is the candor tolerable?

Direct guidance is our position. We need to know where honest becomes discouraging for a real executive.

4

Does the record feel worth keeping?

Continuity is one of our three pillars and we have no evidence for it. Do users come back for the accumulating record, or is it just storage?

Every question already has its evidence

Decided before the build, so three weeks with fifty people produces answers rather than impressions.

Question	Evidence to collect
1 Does the first win land?	Share reaching an editable artifact; time to first artifact; perceived relevance in research sessions
2 Does the loop close?	Artifact-to-used conversion; used-to-outcome conversion; time to outcome update
3 Is candor useful rather than discouraging?	Feedback on narrative summaries and recommendations; abandonment after direct or negative feedback
4 Does the record feel worth keeping?	Return rate; users revisiting prior artifacts or outcomes; stated intent to continue

ALSO TRACKED

Does personalization feel credible on limited inputs?	Share accepting the first recommendation; decline reasons; optional import uptake tracked separately
Does the plan feel focused rather than generic?	Plan adoption; active-action completion; feedback on workload and clarity
Do context-building recommendations help?	Uptake of optional actions, and whether they lead to a completed action or artifact

No measure here depends on a score, a benchmark, or a comparison between users.

Next Steps

These are the decisions the build is waiting on

Each of these blocks something downstream. Dates are the latest they can slip without moving the pilot.

Approve the north star and the five-feature cut	Client	This week
Adopt the Presence Audit as V1’s progress model, and decide whether any score persists internally	Client + D&T	This week
Confirm the direction-not-title onboarding change	Client	This week
Define the minimum enterprise dashboard for V1	Client + D&T	This week
Resolve how org membership is stored without breaking the privacy promise	Client + D&T	This week
Deloitte security and compliance requirements	Client	Immediate, blocking architecture
Terms of service and privacy policy drafted by counsel	Client	Before the November beta
Whether the 50 pilot users must approve to reach December	Client via Deloitte	Before scope lock
Naming and trademark path	Client	Before public launch, not before build

Discovery closes this week

Everything after it passes through one gate: the experience brief.

THIS WEEK

- Revise the features and functionality matrix against this strategy
- Technical session on how recommendations get generated
- Identity-light direction review
- Experience brief: the full requirements document for V1

NEXT

- Design system foundation and component library
- Prototype hub with concepts for the key flows

THEN

- Rolling handoff to engineering, roughly two screens a week
- Pilot cohort recruiting and discussion guides
- First testing sessions against the prototype

Rule number one: nothing gets built without an experience brief

It is how we hold scope on a timeline this compressed, and how we make sure what engineering builds is what we agreed to.

A product earns its next version by working once for real people.

People don't care how much we shipped. They care whether the what they were told to do was worth doing, and whether the product was still useful the second time. That is the only result that matters.